

Identifying Enron Fraud.

A write up of my responses to the questions asked as part of the Udacity's Intro to Machine Learning Project

I have tried to answer these questions in the notebook EnronDataset_EDA.ipynb as well, here just formulating the thought process and steps that are documented in the notebook in brief comments.

Question 1 : Summarize for us the goal of this project and how machine learning is useful in trying to accomplish it. As part of your answer, give some background on the dataset and how it can be used to answer the project question. Were there any outliers in the data when you got it, and how did you handle those? [relevant rubric items: “data exploration”, “outlier investigation”]

In the year 2000, Enron was one of the largest companies in the US. but by 2002, Enron had collapsed as a corporation due to widespread corporate fraud.

In the investigation that followed, huge amount of information was released to public, which included information about executive's salary, stock options, employee emails.

We are given a dataset of 146 data points, each having 21 features ranging from salary, bonus etc to email features like number of messages, POI conversations.

This analysis and modelling once implemented can be used to identify a person as POI (person of interest (POI) is someone who was indicted for fraud, settled with the government, or testified in exchange for immunity) or not..

Following are the major steps in the procedure :

- Exploring Enron Data Set
- Feature Selection, engineering and processing
- Choosing and Running an algorithm
- Validation of results from the algorithm

As part of the project I have done some initial EDA of the dataset, which covers analysis like, distribution of poi to non-poi, relation between salary/bonus, distribution of bonus and salary, relation between total stock value and exercised stock options among others. The EDA can be found in EnronDataset_EDA.ipynb

Out of 146 data points, only 18 are POI, due to complexity of the data set using machine learning to find factors governing and patterns that might lead to a person being categorized 'POI' and can be used to identify POIs in new data set.

Outlier Detection.

On exploring some data points and doing basic visualizations we can see that following observed points can be considered as outliers.

'TOTAL' - This is a total metric which is not be considered for this analysis as it's just consolidation in excel.

'THE TRAVEL AGENCY IN THE PARK' - This is not a person also, most of the attributes for this data point is NaN.

'LOCKHART EUGENE E' is another person who has all of the features as NaN, hence dropping this as well.

Following the data cleaning, 143 data points remained.

Question 2 : What features did you end up using in your POI identifier, and what selection process did you use to pick them? Did you have to do any scaling? Why or why not? As part of the assignment, you should attempt to engineer your own feature that does not come ready-made in the dataset -- explain what feature you tried to make, and the rationale behind it. (You do not necessarily have to use it in the final analysis, only engineer and test it.) In your feature selection step, if you used an algorithm like a decision tree, please also give the feature importances of the features that you use, and if you used an automated feature selection function like SelectKBest, please report the feature scores and reasons for your choice of parameter values. [relevant rubric items: “create new features”, “intelligently select features”, “properly scale features”]

- Features Selection

I did not use any algorithm like kBest, while did feature selection manually.

So based on my EDA, while I couldn't clearly see which features were more important than others, but it was clear there were features in the dataset that huge NaNs, which won't contribute to the ultimate classification.

So I started with eliminating the features that had too many NaN(greater than 75)

Then i dropped email address and person, as they were only there for identification and should play no role in the model.

- Features Creation

I ended up creating 3 features ->

eso/tsv → exercised_Stock_options/total_stock_value

from_poi/to_msg → from_poi_to_this_person / to_messages

to_poi/from_msg → from_this_person_to_poi / from_messages

I did compare the performance with selecting these features with a simple decision tree classifier to see what can be expected and decided to keep all these new features for optimum performance.

- Features Scaling

I did not encounter a need for using feature scaling, if I had used an algo like KMeans, I might have needed to do a max min split scaling.

Question 3 : What algorithm did you end up using? What other one(s) did you try? How did model performance differ between algorithms? [relevant rubric item: “pick an algorithm”]

I tried multiple algorithms like Naive Bayes, RandomForest, SVC, AdaBoost, Bagging.

I found that I was getting best results with RandomForest and AdaBoost and decided to parameter tuning for both of them to see which of them give better metrics.

Details can be seen in the notebook.

Question 4 : What does it mean to tune the parameters of an algorithm, and what can happen if you don't do this well? How did you tune the parameters of your particular algorithm? What parameters did you tune? (Some algorithms do not have parameters that you need to tune -- if this is the case for the one you picked, identify and briefly explain how you would have done it for the model that was not your final choice or a different model that does utilize parameter tuning, e.g. a decision tree classifier). [relevant rubric items: “discuss parameter tuning”, “tune the algorithm”]

Tuning the parameter of the algorithms is essentially changing how the algorithm behaves, this can be done manually with simultaneous plotting, visualizing results and analysing.

Parameter tuning when doing manually is lot of trial and error and kind of an art.

I opted to use sklearn's GridSearchCV for finding the optimum parameters among a range of parameters.

I did parameter tuning for both RandomForest and AdaBoost classifier as I was getting good enough results for both of them.

For AdaBoost, I chose to tune, n_estimators, learning rate and the type of algorithm between 'SAMME.R', 'SAMME'

For Random Forest, I chose to tune - max_depth, min_samples_split, n_estimators, min_samples_leaf and criterion (gini or entropy)

Question 5 : What is validation, and what's a classic mistake you can make if you do it wrong? How did you validate your analysis? [relevant rubric items: "discuss validation", "validation strategy"]

Validation is the process of testing the our model to see if works as expected or not. Typically we divide the data into test,train splits so that we train our model on train data and test on test data.

We can also leverage sklearn's inbuilt function like StratifiedShuffleSplit and StratifiedKFold to validate our model.

I have chosen to use StratifiedShuffleSplit with $n_folds = 10$ to have confidence on the model.

Question 6 : Give at least 2 evaluation metrics and your average performance for each of them. Explain an interpretation of your metrics that says something human-understandable about your algorithm's performance. [relevant rubric item: "usage of evaluation metrics"]

There can be multiple metrics to check the performance of the model. We were required to check for Precision and Recall, I also checked on accuracy to see how model was performing.

Precision - 0.33235 (AdaBoost)

It shows that 33.25% people were correctly classified as 'poi'. This is helpful in minimizing the false positive, as we do not want to have a case in which wrong person was classified as 'poi'. Higher the precision, lesser the number of people being incorrectly classified.

Recall - 0.3385 (AdaBoost)

This means 33.8% of data was retrieved correctly by the model. That is 35.7% was correctly labeled as poi.

Accuracy - 0.82113(AdaBoost)

Model was 82.1% accurate in predicting whether a person is POI or not. Higher the accuracy, better the model.

- List of Resources

- scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html
- <http://www.investopedia.com/updates/enron-scandal-summary/>
- https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.as_matrix.html
- <http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html>
- <https://www.analyticsvidhya.com/blog/2015/11/quick-introduction-boosting-algorithms-machine-learning/>